

KERNEL ARCHITECTURE

BLG413E – System Programming, Practice Session 2

Contents

- System Calls
- Kernel Modules

1-Adding a system call

Requirements (for Ubuntu OS): linux-source, kernel-package, fakeroot, libncurses5-dev

Steps:

- extract linux source
- write new system call
- modify Makefiles
- modify system call table
- modify system call header file
- compile and install new kernel
- reboot to new kernel
- test new system call

Extracting linux source

- move linux source archive file to desktop
 - `cd Desktop`
 - `sudo mv /usr/src/linux-source-3.13.0/linux-source-3.13.0.tar.bz2 linux-source-3.13.0.tar.bz2`
- and extract it
 - `tar -xjvf linux-source-3.13.0.tar.bz2`
- enter linux source folder
 - `cd linux-source-3.13.0`

Writing a system call

- mkdir mycall
- **mycall.c**: under /mycall

```
#include <linux/syscalls.h>
#include <linux/kernel.h>

asmlinkage int sys_mycall(int i, int j){
    return i + j;
}
```

Modifying Makefiles

- create **Makefile** under /mycall
 - write "obj-y := mycall.o" into this file
- modify Makefile under /linux-source-3.13.0 by adding "mycall/" to core-y

```
535 # Objects we will link into vmlinux / subdirs we need to visit
536 init-y      := init/
537 drivers-y   := drivers/ sound/ firmware/ ubuntu/
538 net-y       := net/
539 libs-y      := lib/
540 core-y      := usr/ mycall/
541 endif # KBUILD_EXTMOD
```

Modifying system call table and system call header files

- open arch/x86/syscalls/syscall_32.tbl
 - add "355 i386 mycall sys_mycall" to the end of file

```
363 354 i386 seccomp sys seccomp
364 355 i386 mycall sys_mycall
```

- open include/linux/syscalls.h
 - add "asmlinkage int sys_mycall(int i, int j);" to the end of file before #endif

```
850 asmlinkage long sys_seccomp(unsigned int op, unsigned int flags,
851                             const char __user *uargs);
852 asmlinkage int sys_mycall(int i, int j);
853 #endif
```

Compiling linux kernel

- Configuring the kernel:
 - `make localmodconfig` → include only the modules that are used in the current system
 - You can also use “`make defconfig`” which generates a kernel configuration with the default answer being used for all options
- `make-kpkg clean` → cleans up all from previous kernel compiles
- **Compilation** (**Warning: It may take 1-2 hours**):
`fakeroot`
`make-kpkg --initrd --append-to-version=-custom kernel_image`
`kernel_headers`
- **Output:** two files in parent directory (i.e., Desktop):
 - `linux-image-3.13...deb`
 - `linux-headers-3.13...deb`

Installing compiled kernel

- `sudo dpkg -i linux-image-3.13...`
- `sudo dpkg -i linux-headers-3.13...`
- To make the different kernel versions available at boot time, change the contents of the file `/etc/default/grub` so that the following line is commented out as shown:
 - `#GRUB_HIDDEN_TIMEOUT=0`
- And run the command:
 - `sudo update-grub`
- Then reboot to open from the new kernel (**from Advanced Options for Ubuntu in the boot menu**):
 - `sudo reboot`

Testing new system call

- A simple C program using our new system call to add 2 numbers and print out the result

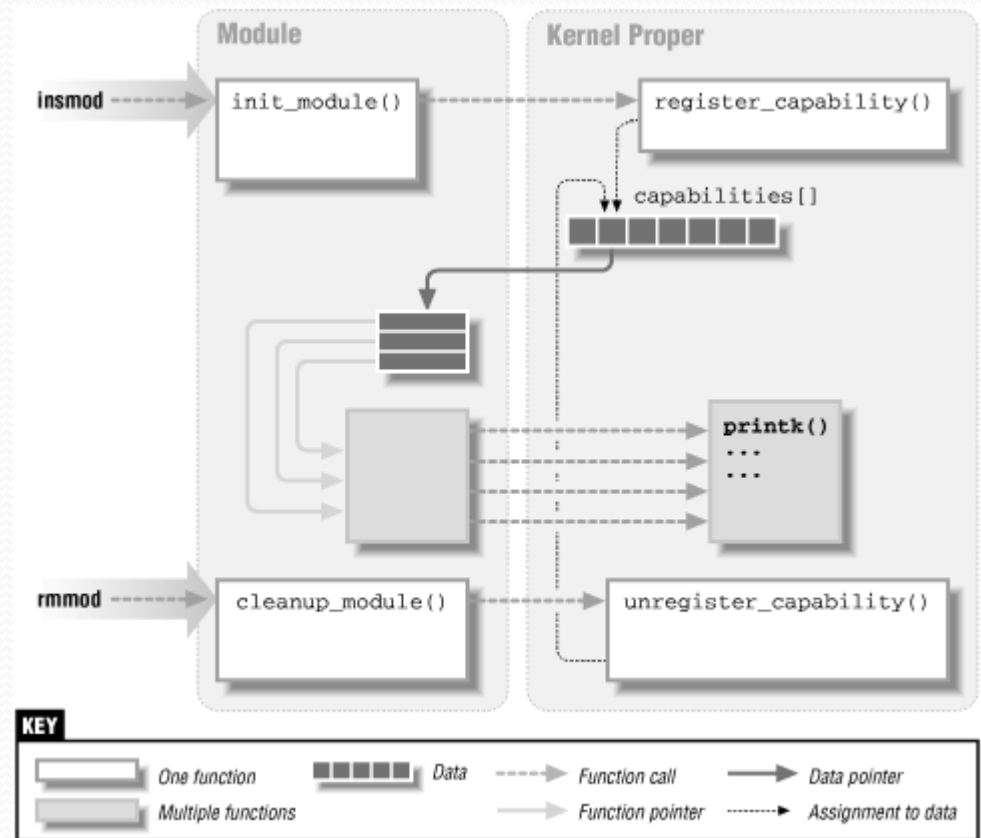
```
1  #include <stdio.h>
2  #define NR_mycall 355
3
4  int main (void){
5      int x1=10, x2=20, y;
6      y = syscall(NR_mycall, x1, x2);
7      printf("%d\n", y);
8      return 0 ;
9  }
10
```

Uninstalling compiled kernel

- When you need to recompile the kernel, **first boot from the original kernel** and uninstall the kernel you have compiled before by using following commands
 - `sudo dpkg -r linux-image-3.13...custom`
 - `sudo dpkg -r linux-headers-3.13...custom`

2-Kernel modules

- A way to add new features to the kernel without rebuilding it.
- Unlike applications, modules register themselves for serving future requests.
- Applications can access the capabilities of a module through system calls.



http://www.xml.com/ldd/chapter/book/figs/ldr2_0201.gif

An example module: hello

- **hello.c:**

```
#include <linux/init.h> /* for module_init and module_exit */
#include <linux/module.h> /* needed by all modules */
MODULE_LICENSE("Dual BSD/GPL"); /* a macro to declare that this module is open source */

static int hello_init(void) /* static: invisible outside the module */
{
    /* to avoid namespace pollution */
    printk(KERN_ALERT "Hello, world\n"); /* printk: kernel print function (macros for priority) */
    return 0; /* KERN_ALERT: a situation requiring immediate action */
}

static void hello_exit(void)
{
    printk(KERN_ALERT "Goodbye, cruel world\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

- **Makefile:**

obj-m := hello.o

all:

make -C /lib/modules/\$(shell uname -r)/build M=\$(PWD) modules

M=\$(PWD) is to build external module in the working directory



Using hello module

- **Compiling:**
 - *make*
- **Loading** (check with ***dmesg*** which is used to see the kernel messages):
 - *sudo insmod ./hello.ko*
- **Unloading** (check with ***dmesg***):
 - *sudo rmmod hello*
- check with ***lsmod*** (which prints the contents of the `/proc/modules` file) before and after loading and unloading

An example module using load time parameters

- **hello.c:**

```
/* $Id: hello.c,v 1.4 2004/09/26 07:02:43 gregkh Exp $ */
#include <linux/init.h>
#include <linux/module.h>
#include <linux/moduleparam.h> /* to enable passing parameters at loadtime */
MODULE_LICENSE("Dual BSD/GPL");

/* A couple of parameters that can be passed in: how many times we say hello, and to whom */
static char *whom = "world";
static int howmany = 1;
module_param(howmany, int, S_IRUGO); /* S_IRUGO: read by the world but cannot be changed */
module_param(whom, charp, S_IRUGO);

static int hello_init(void){
    int i;
    for (i = 0; i < howmany; i++)
        printk(KERN_ALERT "(%d) Hello, %s\n", i, whom);
    return 0;
}

static void hello_exit(void){
    printk(KERN_ALERT "Goodbye, cruel world\n");
}

module_init(hello_init);
module_exit(hello_exit);
```

Specifying module parameters

- `sudo insmod ./hellop.ko whom='Mom' howmany=4`
- `dmesg`

```
4555.764793] (0) Hello, Mom  
4555.764796] (1) Hello, Mom  
4555.764797] (2) Hello, Mom  
4555.764798] (3) Hello, Mom
```

- `sudo rmmod hellop`
- `dmesg`

```
4555.764793] (0) Hello, Mom  
4555.764796] (1) Hello, Mom  
4555.764797] (2) Hello, Mom  
4555.764798] (3) Hello, Mom  
4611.350208] Goodbye, cruel world
```